


Attribute quantisation in the far field: composed against restored



Two separately trained models, and different implementations. Their composed images are nearly indistinguishable and their restorations are not — but this comparison confounds quantisation with trajectory divergence between runs, so it illustrates rather than measures. The within-model comparison, one model in both attribute states, is Figure 4.14 and is not yet rendered.